

Practical ML: Multi-class Classification

Instructor: Marion Neumann

Scribe: Jingyu Xin

Reading:

1 Introduction

Some ML applications/datasets have more than two classes and require us to perform multi-class classification. Some examples are handwritten digit classification (cf. Figure 1), where $\mathcal{C} = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$ or image-based plant disease classification, where the goal is to identify a *specific* disease (cf. Figure 2).

Given a data set $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ where $y_i \in \{1, \dots, k\}$. Which model will you choose to do *multi-class* classification?

Recall: k -NN, Naive Bayes, Decision Trees, and Random Forests, and verify that these models can inherently deal with more than two classes.

Figure 1 (from [LFD CH6.2]¹) illustrates a k -NN solution for the multi-class handwritten digits problem using a simple set of $2d$ features.

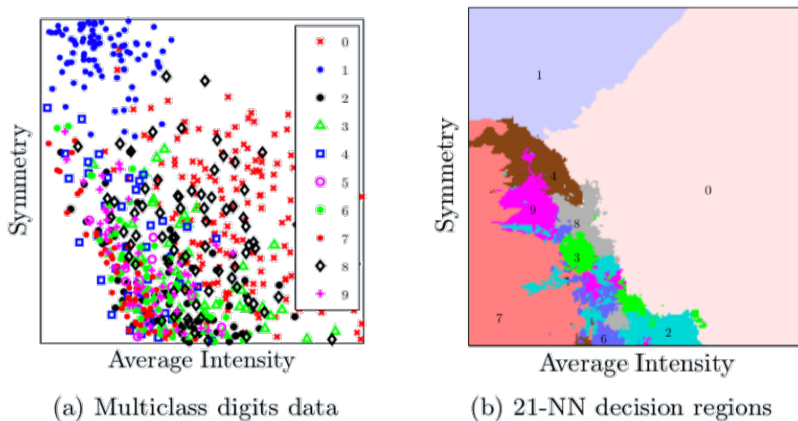


Figure 1: 21-NN decision classifier for the multi-class digits data. [LFD CH6.2]

For **logistic regression** we can simply use the following *conditional likelihood* in a MLE or MAP approach (cf. Section 4 below):

$$p(y = c \mid \mathbf{x}, \mathbf{w}_1, \dots, \mathbf{w}_k) = \frac{e^{\mathbf{w}_c^\top \mathbf{x}}}{\sum_{c=1}^k e^{\mathbf{w}_c^\top \mathbf{x}}}.$$

However, some learning methods do not *easily* extend to multi-class models such as support vector machines (SVMs) and Gaussian processes (GPs). For these cases we can cast multi-class classification as **multiple binary classification problems**.

¹[LFD CH6.2] Learning From Data, Abu-Mostafa et al., 2012, AMLBook

2 1-vs-all Multi-class Classification

Use K classification problems on $\tilde{y}_i = \begin{cases} +1 & \text{if } y_i = c \\ -1 & \text{otherwise} \end{cases}$ for $c \in \{1, \dots, k\}$.

\Rightarrow we get f_1, \dots, f_k models:

$$h(\mathbf{x}) = \arg \max_c f_c(\mathbf{x}) \quad (1)$$

assign the class label of the **most confident** model.

Note:

- only works if classifier **output is comparable** \rightarrow for logistic regression and probabilistic classifiers such as GPs; does not work for SVMs
- for large K , we get class imbalanced

Platt Scaling for SVMs

Since for SVMs the output of different classifiers is not comparable (they are not normalized probabilities), we need to apply a scaling trick to transform the predictions into a probability distribution over classes. This method is commonly referred to as *Platt scaling*.

Let $z_i = f(\mathbf{x}_i) \quad \forall i = 1, \dots, n$ be the SVM predictions before squashing them through the sign function $h(\mathbf{x}) = \text{sign}(f(\mathbf{x})) = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$.

Now, build k datasets $\tilde{D}^{(1)} = \{(z_i, \tilde{y}_i)\}_{i=1}^n, \dots, \tilde{D}^{(k)} = \{(z_i, \tilde{y}_i)\}_{i=1}^n$.

In $\tilde{D}^{(c)}$, $\tilde{y}_i = \begin{cases} +1 & \text{if } y_i = c \\ -1 & \text{otherwise} \end{cases}$ for $c \in \{1, \dots, k\}$. Use those datasets to train k **logistic regression models**

$\tilde{h}_c(z) = \tilde{h}_c(f(\mathbf{x}))$ which rescale the original output $f(\mathbf{x})$. Then, use

$$H(\mathbf{x}) = \arg \max_c \tilde{h}_c(z) \quad (2)$$

3 1-vs-1 Multi-class Classification

Create $\frac{k(k-1)}{2}$ binary classifiers using i vs $j \quad \forall i, j \in \{1, \dots, k\}$.

Assign final label based on (weighted) majority vote. For example:

	1	2	3	4	5
1	N/A	+	+	-	+
2	-	N/A	-	+	+
3	-	+	N/A	-	-
4	+	-	+	N/A	-
5	-	-	+	+	N/A

Table 1: Example of majority vote

The label assigned should be $c = 1$.

Figure 2 illustrates a multi-class plant disease classification approach. The classifier is a one-vs-one multi-class support vector machine (SVM) using a radial basis function (RBF) kernel [NEUMANN, ICPR 2014].²

²[NEUMANN, ICPR 2014] Erosion Band Features for Cell Phone Image Based Plant Disease Classification, M. Neumann, L. Hallau, B. Klatt, K. Kersting, C. Bauckhage, ICPR 2014

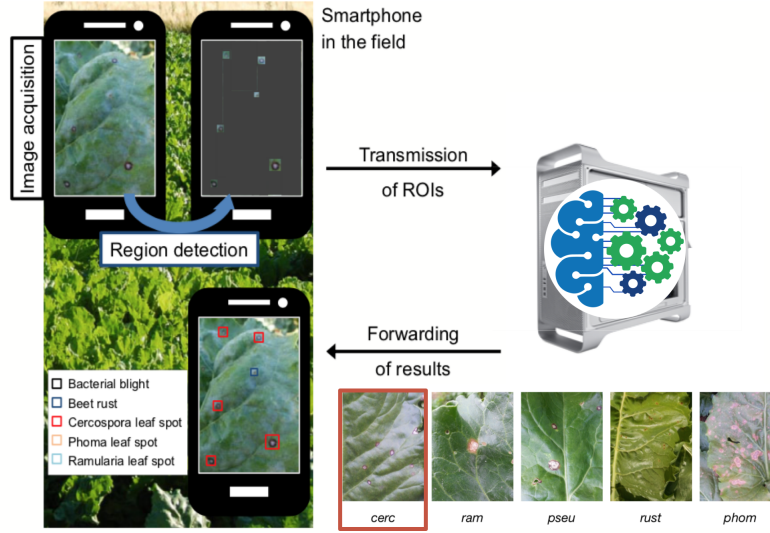


Figure 2: Image-based plant disease classification of sugar beet leaves. Leaf spots can be caused by *Cercospora beticola* (*cerc*), *Ramularia beticola* (*ram*), *Pseudomonas syringae* (*pseu*), *Uromyces betae* (*rust*), or *Phoma betae* (*phom*).

Advantages:

- each problem is class balanced
- each problem is small (few training data)

Disadvantages:

- run time is $\mathcal{O}(k^2)$
- less straight forward to implement

4 [optional] Multi-class Logistic Regression

For logistic regression we can implement 1-vs-all directly as our usual MLE estimation by minimizing the negative log-likelihood. Our prediction model is:

$$p(y = c \mid \mathbf{x}, \mathbf{w}_1, \dots, \mathbf{w}_k) = \frac{e^{\mathbf{w}_c^\top \mathbf{x}}}{\sum_{l=1}^k e^{\mathbf{w}_l^\top \mathbf{x}}} \quad (3)$$

$$h(\mathbf{x}) = \arg \max_c e^{\mathbf{w}_c^\top \mathbf{x}} = \arg \max_c \mathbf{w}_c^\top \mathbf{x}, \quad (4)$$

where the \mathbf{w}_c indicate k hyperplanes.

Training: minimize *negative log-likelihood*:

$$\begin{aligned} \arg \min_W L(W) &= \arg \min_{\mathbf{w}_1, \dots, \mathbf{w}_k} L(\mathbf{w}_1, \dots, \mathbf{w}_k) \\ &= \arg \min_{\mathbf{w}_1, \dots, \mathbf{w}_k} -\log p(\mathbf{y} \mid X, \mathbf{w}_1, \dots, \mathbf{w}_k) \\ &= \arg \min_{\mathbf{w}_1, \dots, \mathbf{w}_k} -\log \prod_{i=1}^n p(y_i \mid \mathbf{x}_i, \mathbf{w}_1, \dots, \mathbf{w}_k) \end{aligned} \quad (5)$$

Note: $p(y \mid \mathbf{x}, \mathbf{w})$ is a multinoulli distribution (one experiment with k categorical outcomes):

$$p(y = c \mid \mathbf{x}, \mathbf{w}) = \prod_{c=1}^k \theta_c^{I(y=c)}.$$

Plugging this into Eq. (5), we get

$$\begin{aligned} \arg \min_W L(W) &= \arg \min_{\mathbf{w}_1, \dots, \mathbf{w}_k} -\log \prod_{i=1}^n \prod_{c=1}^k \left(\frac{e^{\mathbf{w}_c^\top \mathbf{x}_i}}{\sum_{l=1}^k e^{\mathbf{w}_l^\top \mathbf{x}_i}} \right)^{I(y_i=c)} \\ &= \arg \min_{\mathbf{w}_1, \dots, \mathbf{w}_k} \sum_{i=1}^n \left[\log \left(\sum_{l=1}^k e^{\mathbf{w}_l^\top \mathbf{x}_i} \right) - \sum_{c=1}^k I(y_i = c) \mathbf{w}_c^\top \mathbf{x}_i \right] \end{aligned} \quad (6)$$

Now, we can take the derivatives w.r.t. $\mathbf{w}_1, \dots, \mathbf{w}_k$ to get the gradient:

$$g(W) = \sum_{i=1}^n \left(\begin{bmatrix} \mu_{i1} \\ \vdots \\ \mu_{ik} \end{bmatrix} - \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \end{bmatrix} \right) \otimes \mathbf{x}_i \quad (7)$$

where

$$\mathbf{a} \otimes \mathbf{b} = \begin{bmatrix} a_1 \mathbf{b} \\ a_2 \mathbf{b} \\ \vdots \\ a_k \mathbf{b} \end{bmatrix}$$

Just looking at the portion of $g(W)$ corresponding to \mathbf{w}_c , i.e. the derivative of $L(W)$ w.r.t. \mathbf{w}_c , we have:

$$\nabla_{\mathbf{w}_c} L(W) = \sum_i \left(\frac{e^{\mathbf{w}_c^\top \mathbf{x}_i}}{\sum_{l=1}^k e^{\mathbf{w}_l^\top \mathbf{x}_i}} - I(y_i = c) \right) \mathbf{x}_i \quad (8)$$

The Hessian can be derived similarly and both expressions can be plugged into any gradient-based optimizer to get the MLE estimate.